



## **DTI 5126: Fundamentals of Data Science**

### **Assignment 2**

**Prepared by:**  
**Mahmoud Yahia Ahmed**

## Part A:

### 1) Classification:

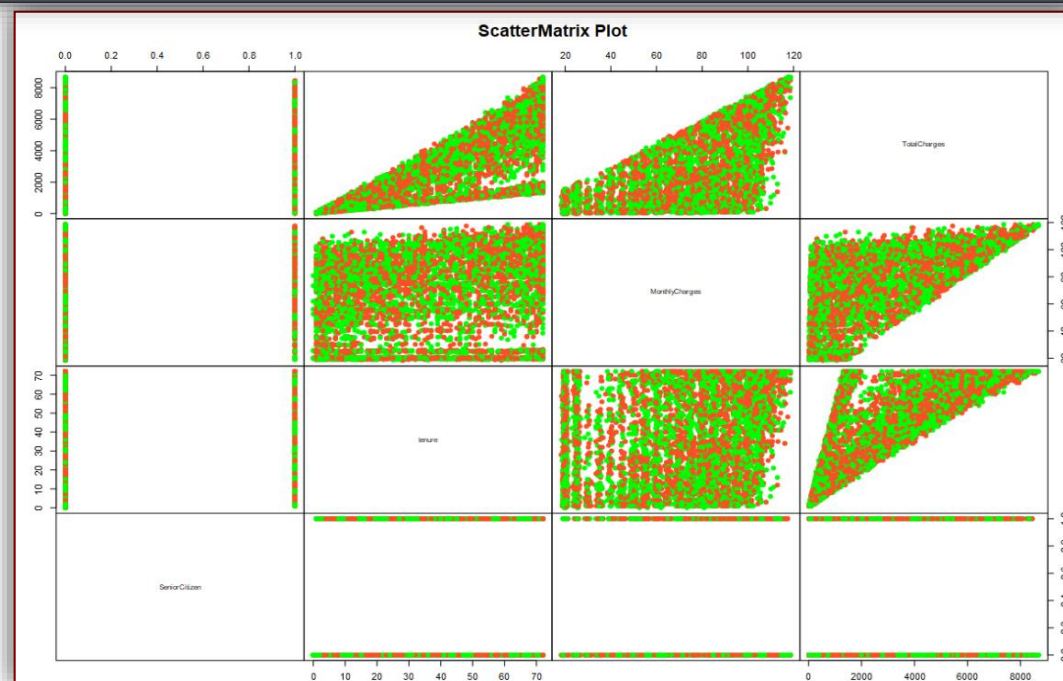
- a) Generate a scatterplot matrix to show the relationships between the variables and a heatmap to determine correlated attributes.

#### I) Reading Data

```
#reading data
DataFile = "C:/Users/mm/Documents/Churn Dataset.csv"
ChurnData = read.csv(DataFile, header=T)
---
```

#### II) Plotting Categorical Data


```
####{r}
##ScatterMatrix Plot for Categorical data
df <- ChurnData[c(3,6,19,20)]
pairs(df, # Data
      pch = 19, # Pch symbol
      col = c("green", "#FC4E27"), # Color
      main = "ScatterMatrix Plot", # Title
      gap = 0, # Subplots distance
      rowlattop = FALSE, # Diagonal direction
      labels = colnames(df), # Labels
      cex.labels = 0.8, # Size of diagonal texts
      font.labels = 1)
---
```



#### III) Plotting Heat map

```
####{r}
###heatmap
dta<- cor(ChurnData[apply(ChurnData,is.numeric)])
cdf <- melt(dta)

ggplot(cdf, aes(Var1, Var2)) +
  geom_tile(aes(fill = value), colour = "white") +
  scale_fill_gradient(low = "white", high = "red")
---
```



**b) Ensure data is in the correct format for downstream processes (e.g., remove redundant information, convert categorical to numerical values, address missing values, etc.).**

### I) Removing Nulls from “TotalCharges”.

```
##{r}
#removing nulls from dataset
sapply(ChurnData, function(x) sum(is.na(x)))
ChurnData <- ChurnData[complete.cases(ChurnData), ]
```

customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService
0	0	0	0	0	0	0
MultipleLines	InternetService	OnlineSecurity	OnlineBackup	DeviceProtection	TechSupport	StreamingTV
0	0	0	0	0	0	0
StreamingMovies	Contract	PaperlessBilling	PaymentMethod	MonthlyCharges	TotalCharges	Churn
0	0	0	0	0	11	0

### II) Replacing Values (“No Internet/Phone service”) to (“No”) in occurrence Columns.

```
##{r}
for(i in 10:15){
  ChurnData[,i] <- as.factor(mapvalues(ChurnData[,i],
                                       from= c("No internet service"), to= c("No")))
}

ChurnData[,8] <- as.factor(mapvalues(ChurnData[,8],
                                       from= c("No phone service"), to= c("No")))
##
```

### III) Since the minimum tenure is 1 month and maximum tenure is 72 months, we can group them into five tenure groups: and putting news groups into new column named(“tenure\_group”)

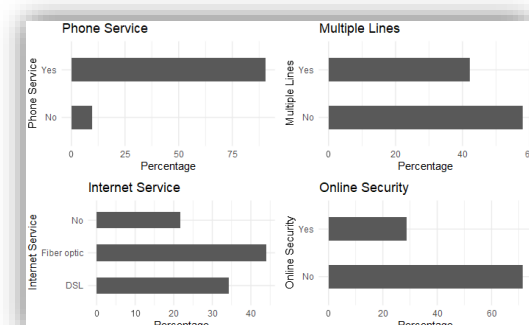
```
##{r}
min(ChurnData$tenure); max(ChurnData$tenure)
group_tenure <- function(tenure){
  if (tenure >= 0 & tenure <= 12){
    return('0-12 Month')
  }else if(tenure > 12 & tenure <= 24){
    return('12-24 Month')
  }else if (tenure > 24 & tenure <= 48){
    return('24-48 Month')
  }else if (tenure > 48 & tenure <=60){
    return('48-60 Month')
  }else if (tenure > 60){
    return('> 60 Month')
  }
}
ChurnData$tenure_group <- sapply(ChurnData$tenure,group_tenure)
ChurnData$tenure_group <- as.factor(ChurnData$tenure_group)
##
```

### IV) Dropping “tenure, Customer ID” columns.

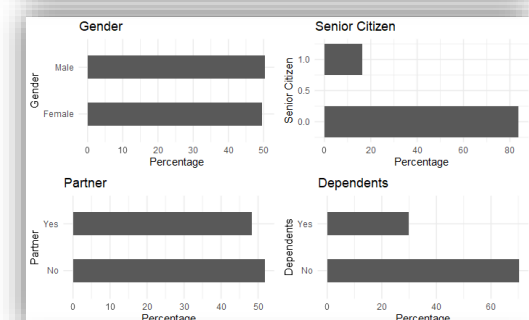
```
##{r}
ChurnData$customerID <- NULL
ChurnData$tenure <- NULL
##
```

### V) Plotting all Categorical Columns to see if it’s possible to drop any unnecessary.


```
##{r}
p5 <- ggplot(ChurnData, aes(x=PhoneService)) + ggtitle("Phone Service") + xlab("Phone Service") +
  geom_bar(aes(y = 100*(..count../sum(..count..)), width = 0.5) + ylab("Percentage") + coord_flip() + theme_minimal())
p6 <- ggplot(ChurnData, aes(x=MultipleLines)) + ggtitle("Multiple Lines") + xlab("Multiple Lines") +
  geom_bar(aes(y = 100*(..count../sum(..count..)), width = 0.5) + ylab("Percentage") + coord_flip() + theme_minimal())
p7 <- ggplot(ChurnData, aes(x=InternetService)) + ggtitle("Internet Service") + xlab("Internet Service") +
  geom_bar(aes(y = 100*(..count../sum(..count..)), width = 0.5) + ylab("Percentage") + coord_flip() + theme_minimal())
p8 <- ggplot(ChurnData, aes(x=OnlineSecurity)) + ggtitle("Online Security") + xlab("Online Security") +
  geom_bar(aes(y = 100*(..count../sum(..count..)), width = 0.5) + ylab("Percentage") + coord_flip() + theme_minimal())
grid.arrange(p5, p6, p7, p8, ncol=2)
```




```
##{r}
p1 <- ggplot(ChurnData, aes(x=gender)) + ggtitle("Gender") + xlab("Gender") +
  geom_bar(aes(y = 100*(..count../sum(..count..)), width = 0.5) + ylab("Percentage") + coord_flip() + theme_minimal())
p2 <- ggplot(ChurnData, aes(x=SeniorCitizen)) + ggtitle("Senior Citizen") + xlab("Senior Citizen") +
  geom_bar(aes(y = 100*(..count../sum(..count..)), width = 0.5) + ylab("Percentage") + coord_flip() + theme_minimal())
p3 <- ggplot(ChurnData, aes(x=Partner)) + ggtitle("Partner") + xlab("Partner") +
  geom_bar(aes(y = 100*(..count../sum(..count..)), width = 0.5) + ylab("Percentage") + coord_flip() + theme_minimal())
p4 <- ggplot(ChurnData, aes(x=Dependents)) + ggtitle("Dependents") + xlab("Dependents") +
  geom_bar(aes(y = 100*(..count../sum(..count..)), width = 0.5) + ylab("Percentage") + coord_flip() + theme_minimal())
grid.arrange(p1, p2, p3, p4, ncol=2)
```



```
##{r}
p9 <- ggplot(ChurnData, aes(x=OnlineBackup)) + ggtitle("Online Backup") + xlab("Online Backup") +
  geom_bar(aes(y = 100*(..count../sum(..count..)), width = 0.5) + ylab("Percentage") + coord_flip() + theme_minimal())
p10 <- ggplot(ChurnData, aes(x=DeviceProtection)) + ggtitle("Device Protection") + xlab("Device Protection") +
  geom_bar(aes(y = 100*(..count../sum(..count..)), width = 0.5) + ylab("Percentage") + coord_flip() + theme_minimal())
p11 <- ggplot(ChurnData, aes(x=TechSupport)) + ggtitle("Tech Support") + xlab("Tech Support") +
  geom_bar(aes(y = 100*(..count../sum(..count..)), width = 0.5) + ylab("Percentage") + coord_flip() + theme_minimal())
p12 <- ggplot(ChurnData, aes(x=StreamingTV)) + ggtitle("Streaming TV") + xlab("Streaming TV") +
  geom_bar(aes(y = 100*(..count../sum(..count..)), width = 0.5) + ylab("Percentage") + coord_flip() + theme_minimal())
grid.arrange(p9, p10, p11, p12, ncol=2)
```



```
##{r}
p13 <- ggplot(ChurnData, aes(x=StreamingMovies)) + ggtitle("Streaming Movies") + xlab("Streaming Movies") +
  geom_bar(aes(y = 100*(..count../sum(..count..)), width = 0.5) + ylab("Percentage") + coord_flip() + theme_minimal())
p14 <- ggplot(ChurnData, aes(x=Contract)) + ggtitle("Contract") + xlab("Contract") +
  geom_bar(aes(y = 100*(..count../sum(..count..)), width = 0.5) + ylab("Percentage") + coord_flip() + theme_minimal())
p15 <- ggplot(ChurnData, aes(x=PaperlessBilling)) + ggtitle("Paperless Billing") + xlab("Paperless Billing") +
  geom_bar(aes(y = 100*(..count../sum(..count..)), width = 0.5) + ylab("Percentage") + coord_flip() + theme_minimal())
p16 <- ggplot(ChurnData, aes(x=PaymentMethod)) + ggtitle("Payment Method") + xlab("Payment Method") +
  geom_bar(aes(y = 100*(..count../sum(..count..)), width = 0.5) + ylab("Percentage") + coord_flip() + theme_minimal())
p17 <- ggplot(ChurnData, aes(x=tenure_group)) + ggtitle("Tenure Group") + xlab("Tenure Group") +
  geom_bar(aes(y = 100*(..count../sum(..count..)), width = 0.5) + ylab("Percentage") + coord_flip() + theme_minimal())
grid.arrange(p13, p14, p15, p16, p17, ncol=2)
```



All of the categorical variables seem to have a reasonably broad distribution, therefore, all of them will be kept for the further analysis.

## VI) Convert categorical to numerical values.

```
{r}
columns <- c("Partner", "Dependents", "PhoneService",
             "MultipleLines", "OnlineSecurity", "OnlineBackup",
             "DeviceProtection", "TechSupport", "StreamingTV",
             "StreamingMovies", "PaperlessBilling", "Churn", "gender",
             "InternetService", "Contract", "PaymentMethod", "tenure_group")
for (col in names(ChurnData)){
  for (j in (columns)){
    if(ChurnData[col]%in%ChurnData[j]){
      ChurnData[,j] = as.numeric(as.factor(ChurnData[,col])) - 1
    }
  }
}
```

## VII) Checking for the highest correlated feature and removing it.

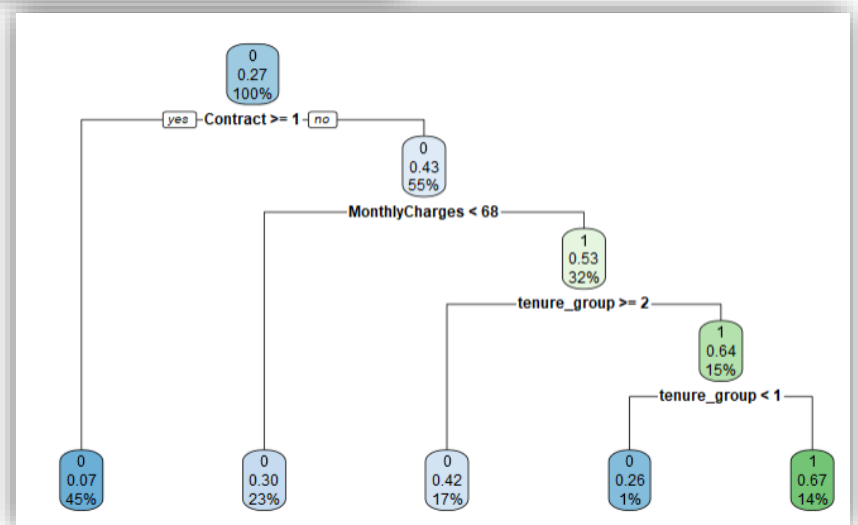
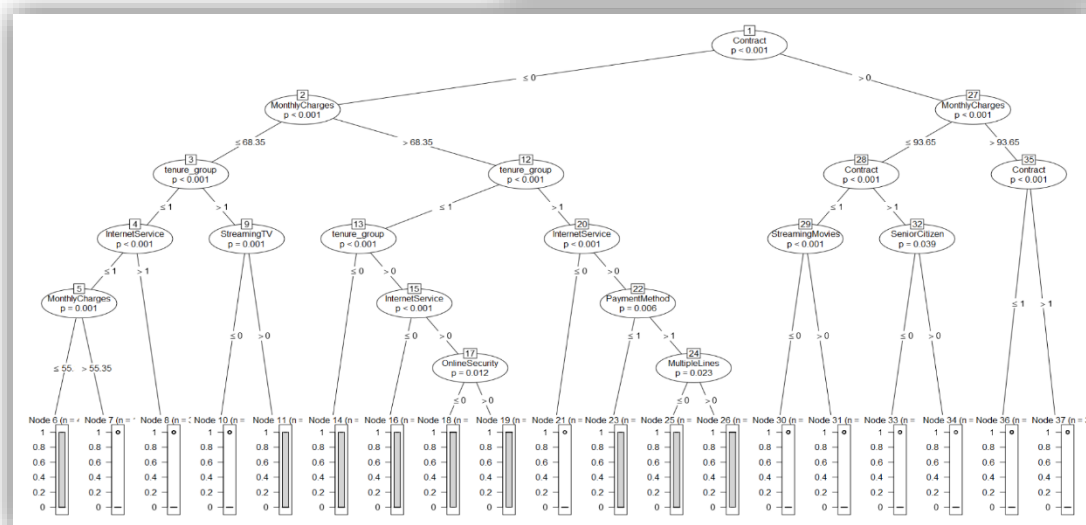
```
{r}

correlationMatrix <- cor(ChurnData)
highlyCorrelated <- findCorrelation(correlationMatrix, cutoff=0.65)
# print indexes of highly correlated attributes
print(highlyCorrelated)
hc = sort(highlyCorrelated)|
```

```
{r}
ChurnData$TotalCharges<- NULL
```

## c) Split the dataset into 80 training/20 test set and fit a decision tree to the training data. Plot the tree, and interpret the results.

```
{r}
set.seed(1112)
split_train_test <- createDataPartition(ChurnData$Churn,p=0.8,list=FALSE)
dtrain<- ChurnData[split_train_test,]
dtest<- ChurnData[-split_train_test,]
#Plot decision tree
rtree <- rpart(Churn ~ ., dtrain ,method ="class")
#plot conditional parting plot
ctree_ <- ctree(Churn ~ ., dtrain)
plot(ctree_)
rpart.plot(rtree)
pred <- predict(rtree , newdata = dtest,type="class")
confusionMatrix(pred, as.factor( dtest$Churn) ) #check accuracy
```



```
Reference
Prediction  0  1
0    1004  218
1     49   135
Accuracy : 0.8101
95% CI : (0.7886, 0.8303)
No Information Rate : 0.7489
P-Value [Acc > NIR] : 3.069e-08
Kappa : 0.3995
McNemar's Test P-Value : < 2.2e-16
Sensitivity : 0.9535
Specificity : 0.3824
Pos Pred Value : 0.8216
Neg Pred Value : 0.7337
Precision : 0.8216
Recall : 0.9535
F1 : 0.8826
Prevalence : 0.7489
Detection Rate : 0.7141
Detection Prevalence : 0.8691
Balanced Accuracy : 0.6680
'Positive' Class : 0
```

### From the shown confusion matrix

The model predicts most of the test points right but there is something confusing with the model so it misses some of the points



- d) Try different ways to improve the decision tree algorithm (e.g., use different splitting strategies, prune the tree after splitting).

```

{r}

ethnancetree2 <- rpart(Churn ~ ., dtrain, method = "class", parms = list(split = "information"))
epred2 <- predict(ethnancetree2, newdata = dtest, type = "class", mode = "everything")
confusionMatrix(epred2, as.factor(dtest$Churn)) #check accuracy

```

```

{r}

DT_poprune <- rpart(Churn ~ ., dtrain,
  method = "class",
  control = rpart.control(cp = 0.0005) )
plotcp(DT_poprune)
DT_model_pruned <- prune(DT_poprune, cp = 0.004)
Pred_poprune <- predict(DT_model_pruned, newdata = subset(dtest, select = -c(Churn)), type = "class")
table_mat_prune <- table(dtest$Churn, Pred_poprune)
acc_Test_prune <- sum(diag(table_mat_prune)) / sum(table_mat_prune)
print(paste("Accuracy:", acc_Test_prune*100))

confusionMatrix(Pred_poprune, as.factor(dtest$Churn), mode = "everything") #check accuracy

```

```

Confusion Matrix and Statistics

      Reference
Prediction 0  1
0    1004  218
1      49  135

      Accuracy : 0.8101
      95% CI   : (0.7886, 0.8303)
No Information Rate : 0.7489
P-Value [Acc > NIR] : 3.069e-08

      Kappa : 0.3995

McNemar's Test P-Value : < 2.2e-16

      Sensitivity : 0.9535
      Specificity : 0.3824
      Pos Pred Value : 0.8216
      Neg Pred Value : 0.7337
      Precision : 0.8216
      Recall : 0.9535
      F1 : 0.8826
      Prevalence : 0.7489
      Detection Rate : 0.7141
      Detection Prevalence : 0.8691
      Balanced Accuracy : 0.6680

```

```

Reference
Prediction 0  1
0    976 187
1     77 166

      Accuracy : 0.8122
      95% CI   : (0.7908, 0.8323)
No Information Rate : 0.7489
P-Value [Acc > NIR] : 1.002e-08

      Kappa : 0.443

McNemar's Test P-Value : 1.967e-11

      Sensitivity : 0.9269
      Specificity : 0.4703
      Pos Pred Value : 0.8392
      Neg Pred Value : 0.6831
      Precision : 0.8392
      Recall : 0.9269
      F1 : 0.8809
      Prevalence : 0.7489
      Detection Rate : 0.6942
      Detection Prevalence : 0.8272
      Balanced Accuracy : 0.6986

'Positive' Class : 0

```

Does pruning the tree improve the accuracy? Yes, but not that much

- e) Classify the data using the XGBoost model with nrounds = 70 and max depth = 3. Evaluate the performance.

```

{r}
#splitting for XGBoost classifier

xtrain = data.matrix(dtrain[1:5626, -18])
ytrain = data.matrix(dtrain[1:5626, 18])
xgtest = data.matrix(dtest[1:5626, -18])
ygtest = data.matrix(dtest[1:5626, 18])

# train an xgboost model
bstDense <- xgboost(data = xtrain, label = ytrain,
  max.depth = 3, eta = 0.2,
  nrounds = 70, gamma = 0.8)

```

```

Confusion Matrix and Statistics

      Reference
Prediction 0  1
0    978 175
1     75 178

      Accuracy : 0.8222
      95% CI   : (0.8012, 0.8418)
No Information Rate : 0.7489
P-Value [Acc > NIR] : 3.116e-11

      Kappa : 0.478

McNemar's Test P-Value : 3.818e-10

      Sensitivity : 0.9288
      Specificity : 0.5042
      Pos Pred Value : 0.8482
      Neg Pred Value : 0.7036
      Precision : 0.8482
      Recall : 0.9288
      F1 : 0.8867
      Prevalence : 0.7489
      Detection Rate : 0.6956
      Detection Prevalence : 0.8201
      Balanced Accuracy : 0.7165


'Positive' Class : 0

```

Is there any sign of overfitting? **No**

**f) Train a deep neural network using Keras with 3 dense layers. Try changing the activation function or dropout rate.**

**I) Using Relu activation function and softmax in the end.**



```

Reference
Prediction 0 1
0 930 123
1 138 215

Accuracy : 0.8144
95% CI : (0.793, 0.8344)
No Information Rate : 0.7596
P-Value [Acc > NIR] : 4.602e-07


Kappa : 0.4993

McNemar's Test P-Value : 0.3862

Sensitivity : 0.8708
Specificity : 0.6361
Pos Pred Value : 0.8832
Neg Pred Value : 0.6091
Precision : 0.8832
Recall : 0.8708
F1 : 0.8769
Prevalence : 0.7596
Detection Rate : 0.6615
Detection Prevalence : 0.7489
Balanced Accuracy : 0.7534

'Positive' Class : 0
    
```

**II) Using (Tanh) activation function and softmax in the end.**



```

Reference
Prediction 0 1
0 920 133
1 132 221

Accuracy : 0.8115
95% CI : (0.7901, 0.8317)
No Information Rate : 0.7482
P-Value [Acc > NIR] : 1.038e-08

Kappa : 0.4993

McNemar's Test P-Value : 1

Sensitivity : 0.8745
Specificity : 0.6243
Pos Pred Value : 0.8737
Neg Pred Value : 0.6261
Precision : 0.8737
Recall : 0.8745
F1 : 0.8741
Prevalence : 0.7482
Detection Rate : 0.6543
Detection Prevalence : 0.7489
Balanced Accuracy : 0.7494

'Positive' Class : 0
    
```

What effects does any of these have on the result?

**No effects when I replaced the Relu function obtained higher accuracy**

**g) Compare the performance of the models in terms of the following criteria: precision, recall, accuracy, F-measure. Identify the model that performed best and worst according to each criterion.**

**DNN model(RELU) achieved the Highest Precision among all models**

the classifier is very conservative - does not risk too much in saying that a sample is Positive

**Base Decisiontree&enhanced achieved the Highest recall among all models**

The model maximizes the number of True Positives But it could be wrong sometimes!

**Pruned Decision tree achieved the Highest F1 among all models**

high F-score can be the result of an imbalance between Precision and Recall

DT&enhanced DT **worst model**

Pruned DT

XGBOOST classifier

Dnn relu AF **best model**

Dnn tanh AF

```

Reference
Prediction 0 1
0 1004 218
1 49 135

Accuracy : 0.8101
95% CI : (0.7886, 0.8303)
No Information Rate : 0.7489
P-Value [Acc > NIR] : 3.069e-08

Kappa : 0.3995

McNemar's Test P-Value : < 2.2e-16

Sensitivity : 0.9535
Specificity : 0.3824
Pos Pred Value : 0.8216
Neg Pred Value : 0.7337
Precision : 0.8216
Recall : 0.9535
F1 : 0.8826
Prevalence : 0.7489
Detection Rate : 0.7141
Detection Prevalence : 0.8691
Balanced Accuracy : 0.6680

'Positive' Class : 0
    
```

```

Reference
Prediction 0 1
0 976 187
1 77 166

Accuracy : 0.8122
95% CI : (0.7908, 0.8323)
No Information Rate : 0.7489
P-Value [Acc > NIR] : 1.002e-08

Kappa : 0.443

McNemar's Test P-Value : 1.967e-11

Sensitivity : 0.9269
Specificity : 0.4703
Pos Pred Value : 0.8392
Neg Pred Value : 0.6831
Precision : 0.8392
Recall : 0.9269
F1 : 0.8809
Prevalence : 0.7489
Detection Rate : 0.6942
Detection Prevalence : 0.8272
Balanced Accuracy : 0.6986

'Positive' Class : 0
    
```

```

Confusion Matrix and Statistics
Reference
Prediction 0 1
0 978 175
1 75 178

Accuracy : 0.8222
95% CI : (0.8012, 0.8418)
No Information Rate : 0.7489
P-Value [Acc > NIR] : 3.116e-11

Kappa : 0.478

McNemar's Test P-Value : 3.818e-10

Sensitivity : 0.9288
Specificity : 0.5042
Pos Pred Value : 0.8482
Neg Pred Value : 0.7036
Precision : 0.8482
Recall : 0.9288
F1 : 0.8867
Prevalence : 0.7489
Detection Rate : 0.6956
Detection Prevalence : 0.8201
Balanced Accuracy : 0.7165

'Positive' Class : 0
    
```

```

Reference
Prediction 0 1
0 930 123
1 138 215

Accuracy : 0.8144
95% CI : (0.793, 0.8344)
No Information Rate : 0.7596
P-Value [Acc > NIR] : 4.602e-07

Kappa : 0.4993

McNemar's Test P-Value : 0.3862

Sensitivity : 0.8708
Specificity : 0.6361
Pos Pred Value : 0.8832
Neg Pred Value : 0.6091
Precision : 0.8832
Recall : 0.8708
F1 : 0.8769
Prevalence : 0.7596
Detection Rate : 0.6615
Detection Prevalence : 0.7489
Balanced Accuracy : 0.7534

'Positive' Class : 0
    
```

```

Reference
Prediction 0 1
0 920 133
1 132 221

Accuracy : 0.8115
95% CI : (0.7901, 0.8317)
No Information Rate : 0.7482
P-Value [Acc > NIR] : 1.038e-08

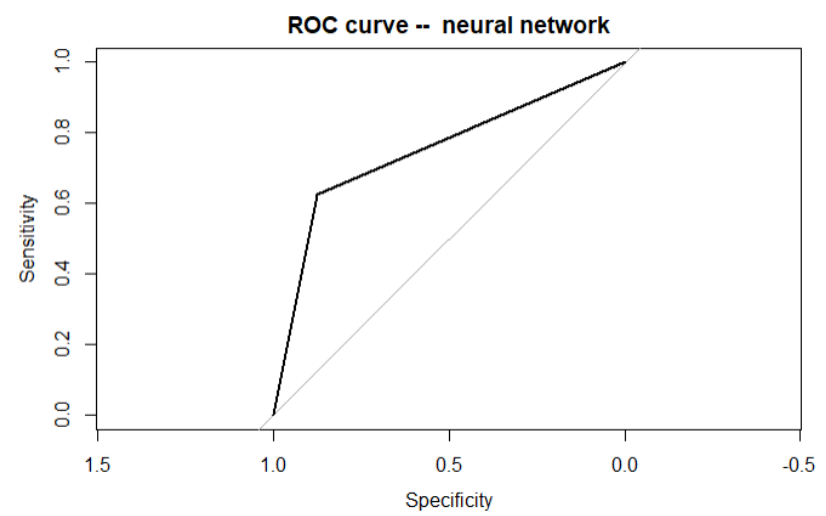

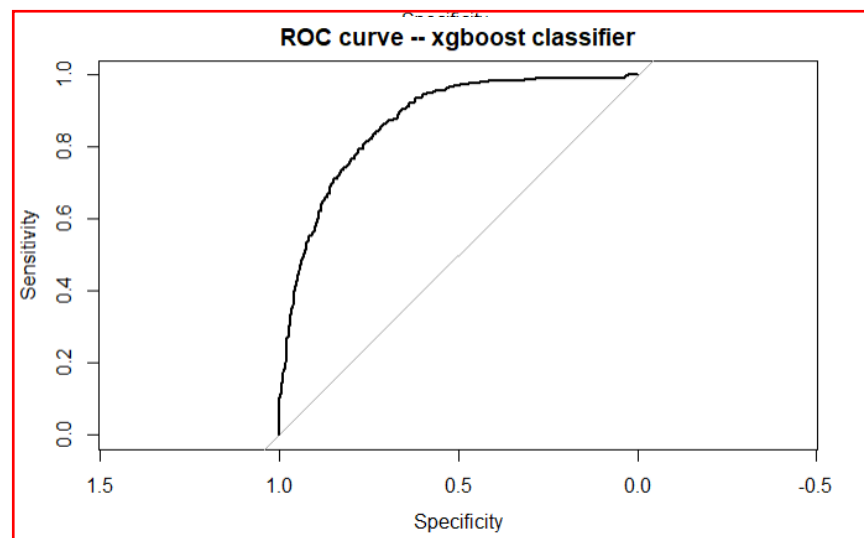
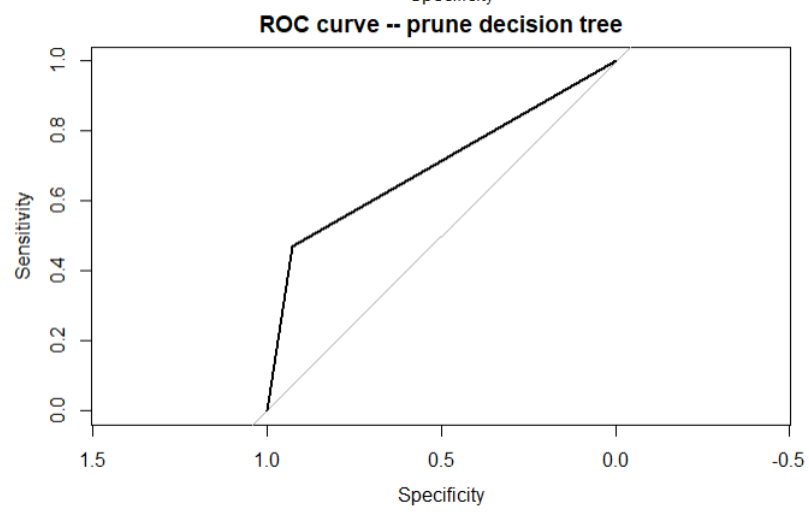


Kappa : 0.4993

McNemar's Test P-Value : 1

Sensitivity : 0.8745
Specificity : 0.6243
Pos Pred Value : 0.8737
Neg Pred Value : 0.6261
Precision : 0.8737
Recall : 0.8745
F1 : 0.8741
Prevalence : 0.7482
Detection Rate : 0.6543
Detection Prevalence : 0.7489
Balanced Accuracy : 0.7494

'Positive' Class : 0
    
```

**h) Use a ROC graph to compare the performance of the DT, XGboost & DNN techniques.**





**The highlighted curve has the high Area under the curve so It has better performance**

## Part B:

**1) Generate a plot of the top 10 transactions.**

```
{r}
#reading data
df = "C:/Users/Queen/Documents/transactions.csv"
trans = read.transactions(df, sep = ",")
#plotting most 10 frq transacions
itemFrequencyPlot(trans, topN = 10)
image(trans[1:10])
```



2) Generate association rules using minimum support of 0.002, minimum confidence of 0.20, and maximum length of 3. Display the rules, sorted by descending lift value.

```
```{r}
# default settings result in zero rules learned
apriori(trans)

transrules ← apriori(trans, parameter = list(support = 0.002, confidence = 0.20, maxlen = 3))
lift ← inspect(sort(transrules, by="lift"))

```
```

|        | lhs<br><chr>                  |    | rhs<br><chr>    | support<br><dbl> | confidence<br><dbl> |  |
|--------|-------------------------------|----|-----------------|------------------|---------------------|--|
| [991]  | {milk, mineral water}         | => | {spaghetti}     | 0.015731236      | 0.3277778           |  |
| [992]  | {milk, spaghetti}             | => | {chocolate}     | 0.010931876      | 0.3082707           |  |
| [993]  | {burgers, french fries}       | => | {green tea}     | 0.005465938      | 0.2484848           |  |
| [994]  | {olive oil, spaghetti}        | => | {chocolate}     | 0.007065725      | 0.3081395           |  |
| [995]  | {almonds}                     | => | {green tea}     | 0.005065991      | 0.2483660           |  |
| [996]  | {mineral water, oil}          | => | {spaghetti}     | 0.002399680      | 0.3272727           |  |
| [997]  | {olive oil, spaghetti}        | => | {mineral water} | 0.010265298      | 0.4476744           |  |
| [998]  | {french fries, grated cheese} | => | {chocolate}     | 0.003199573      | 0.3076923           |  |
| [999]  | {soup, tomatoes}              | => | {spaghetti}     | 0.002266364      | 0.3269231           |  |
| [1000] | {protein bar, spaghetti}      | => | {mineral water} | 0.002266364      | 0.4473684           |  |

991-1000 of 2,023 rows | 1-6 of 8 columns

Previous1...9596979899100Next

3) Select the rule from QII-b with the greatest lift. Compare this rule with the highest lift rule for maximum length of 2.

```
```{r}
# set better support and confidence levels to learn more rules
transrules1 ← apriori(trans, parameter = list(support = 0.002, confidence = 0.20, maxlen = 2))
lift ← inspect(sort(transrules1, by="lift"))
transrules1

```
```

|       | lhs<br><chr>      |    | rhs<br><chr>    | support<br><dbl> | confidence<br><dbl> | coverage<br><dbl> | lift<br><dbl> | count<br><int> |
|-------|-------------------|----|-----------------|------------------|---------------------|-------------------|---------------|----------------|
| [351] | {green tea}       | => | {mineral water} | 0.031062525      | 0.2351160           | 0.132115718       | 0.9863565     | 233            |
| [352] | {parmesan cheese} | => | {mineral water} | 0.004666045      | 0.2348993           | 0.019864018       | 0.9854474     | 35             |
| [353] | {melons}          | => | {mineral water} | 0.002799627      | 0.2333333           | 0.011998400       | 0.9788777     | 21             |
| [354] | {candy bars}      | => | {mineral water} | 0.002266364      | 0.2328767           | 0.009732036       | 0.9769621     | 17             |
| [355] | {light mayo}      | => | {mineral water} | 0.006265831      | 0.2303922           | 0.027196374       | 0.9665389     | 47             |
| [356] | {escalope}        | => | {mineral water} | 0.017064391      | 0.2151261           | 0.079322757       | 0.9024947     | 128            |
| [357] | {energy drink}    | => | {mineral water} | 0.005599253      | 0.2100000           | 0.026663112       | 0.8809899     | 42             |
| [358] | {yogurt cake}     | => | {mineral water} | 0.005599253      | 0.2048780           | 0.027329689       | 0.8595024     | 42             |

351-358 of 358 rows

Previous1...313233343536Next

I) Which rule has the better lift? Which rule has the greater support?

The second rule has the greater support

II) If you were a marketing manager, and could fund only one of these rules, which would it be, and why?

The second rule, because it has the highest support and trust, it will reflect positively on my sales!